

# CLUSTERING

---

## BIRCH Algorithm

# Clustering Large Datasets

- Clustering algorithms we've seen are not good for large datasets because they
  - Assume dataset and data structures used fit in main memory
  - Scan the dataset many times
  - Assume the whole dataset is present at once
    - An **incremental** algorithm can handle dynamic and incoming data
    - An **online** algorithm is capable of providing the best answer so far

# BIRCH

- BIRCH stands for Balanced Iterative Reducing and Clustering using Hierarchies
- It can handle large datasets by using data structures for summarizing information about the clusters
- Scalable and incremental

# Clustering Feature (CF)

- Definition: A **clustering feature (CF)** is a triple (N, **LS**, **SS**) summarizing information about clusters or sub-clusters of objects where:
  - N is the number of objects
  - **LS** is the linear sum ( $\sum_{i=1}^N \mathbf{t}_i$ )                      sum the individual coordinates
  - **SS** is square sum ( $\sum_{i=1}^N \mathbf{t}_i^2$ )                      sum the squared coordinates

# Clustering Feature Example

- Find the CF for a cluster that contains the following data points:

$$t1 = (3, 4), t2 = (2, 6), t3 = (4, 5), t4 = (4, 7), \text{ and } t5 = (3, 8)$$

- CF = (N, **LS**, **SS**)

- N = 5

- $LS = \sum_{i=1}^N t_i = t1 + t2 + t3 + t4 + t5$

- $LS = (3, 4) + (2, 6) + (4, 5) + (4, 7) + (3, 8)$

$$= (3 + 2 + 4 + 4 + 3, 4 + 6 + 5 + 7 + 8) = (16, 30)$$

- $SS = \sum_{i=1}^N t_i^2 = (3^2 + 2^2 + 4^2 + 4^2 + 3^2, 4^2 + 6^2 + 5^2 + 7^2 + 8^2)$

- $SS = (9 + 4 + 16 + 16 + 9, 16 + 36 + 25 + 49 + 64) = (54, 190)$

- CF = (N, **LS**, **SS**) = (5, (16, 30), (54, 190))

# CF Additivity Property

- Let  $CF_1 = (N_1, \mathbf{LS}_1, \mathbf{SS}_1)$  and  $CF_2 = (N_2, \mathbf{LS}_2, \mathbf{SS}_2)$  be the CF vectors of two **disjoint** clusters  $K_1$  and  $K_2$ .

The CF vector of the cluster that is formed by merging  $K_1$  and  $K_2$  is:


$$CF_1 + CF_2 = (N_1 + N_2, \mathbf{LS}_1 + \mathbf{LS}_2, \mathbf{SS}_1 + \mathbf{SS}_2)$$

- Proof: straight forward algebra

# CF Additivity Property - Example

- Let  $CF_1 = (3, (4, 4), (6, 6))$  and  $CF_2 = (3, (11, 5), (41, 9))$  be the CF vectors of two disjoint clusters  $K_1$  and  $K_2$ . Find the CF vector for the cluster that results from merging the two clusters.

$$\begin{aligned} CF &= CF_1 + CF_2 = (N_1 + N_2, \mathbf{LS}_1 + \mathbf{LS}_2, \mathbf{SS}_1 + \mathbf{SS}_2) \\ &= (3 + 3, (4 + 11, 4 + 5), (6 + 41, 6 + 9)) \\ &= (6, (15, 9), (47, 15)) \end{aligned}$$



Cluster 1 =  $\{(1, 1), (2, 1), (1, 2)\}$   
Cluster 2 =  $\{(3, 2), (4, 1), (4, 2)\}$

# Using the CF to Compute Cluster Parameters

- Knowing the CF vector of a cluster, it is easy to compute the cluster's centroid, radius, and diameter
- Remember, for a cluster  $K_m = \{t_{m1}, t_{m2}, \dots, t_{mN}\}$  of  $N$  points:

$$\textit{centroid} = C_m = \frac{\sum_{i=1}^N (t_{mi})}{N}$$

$$\textit{radius} = R_m = \sqrt{\frac{\sum_{i=1}^N (t_{mi} - C_m)^2}{N}}$$

$$\textit{diameter} = D_m = \sqrt{\frac{\sum_{i=1}^N \sum_{j=1}^N (t_{mi} - t_{mj})^2}{(N)(N-1)}}$$



# Using the CF to Compute Cluster Parameters – Centroid

- Knowing the CF vector of a cluster, it is easy to compute the cluster's centroid, radius, and diameter
- For a cluster  $K_m = \{t_{m1}, t_{m2}, \dots, t_{mN}\}$  of  $N$  points:

$$\textit{centroid} = C_m = \frac{\sum_{i=1}^N (t_{mi})}{N} = \frac{LS}{N}$$

# Using the CF to Compute Cluster Parameters – Radius

- Knowing the CF vector of a cluster, it is easy to compute the cluster's centroid, radius, and diameter
- For a cluster  $K_m = \{t_{m1}, t_{m2}, \dots, t_{mN}\}$  of  $N$  points:

$$radius = R_m = \sqrt{\frac{\sum_{i=1}^N (t_{mi} - C_m)^2}{N}}$$

$$R_m = \sqrt{\frac{\sum_{i=1}^N (t_{mi}^2 - 2t_{mi}C_m + C_m^2)}{N}} = \sqrt{\frac{\sum_{i=1}^N t_{mi}^2 - 2C_m \sum_{i=1}^N t_{mi} + \sum_{i=1}^N C_m^2}{N}}$$

$$R_m = \sqrt{\frac{SS - 2C_m LS + NC_m^2}{N}}$$

# Using the CF to Compute Cluster Parameters – Diameter

- Knowing the CF vector of a cluster, it is easy to compute the cluster's centroid, radius, and diameter
- For a cluster  $K_m = \{t_{m1}, t_{m2}, \dots, t_{mN}\}$  of  $N$  points:

$$\text{diameter} = D_m = \sqrt{\frac{\sum_{i=1}^N \sum_{j=1}^N (t_{mi} - t_{mj})^2}{(N)(N-1)}}$$

We can show,  $D_m = \sqrt{\frac{2NSS - 2LS^2}{N(N-1)}}$

# Using the CF to Compute Distance Between Two Cluster

- The centroid Euclidean distance  $D_2$  between two clusters  $K_1$  and  $K_2$  is

$$\begin{aligned} D_2 &= \sqrt{(C_1 - C_2)^2} \\ &= \sqrt{\left(\frac{LS1}{N1} - \frac{LS2}{N2}\right)^2} \end{aligned}$$

- The centroid Manhattan distance  $D_1$  between two clusters  $K_1$  and  $K_2$  is


$$\begin{aligned} D_1 &= |C_1 - C_2| = \sum_{i=1}^d |C_{1i} - C_{2i}|, \text{ where } d \text{ is number of dimensions} \\ &= \left| \frac{LS1}{N1} - \frac{LS2}{N2} \right| \end{aligned}$$

# Example

- Let  $CF_1 = (3, (4, 4), (6, 6))$  and  $CF_2 = (3, (11, 5), (41, 9))$  be the CF vectors of two clusters  $K_1$  and  $K_2$ .

- $$D_2 = \sqrt{(C_1 - C_2)^2} = \sqrt{\left(\frac{LS1}{N1} - \frac{LS2}{N2}\right)^2}$$
$$= \sqrt{\left(\frac{(4,4)}{3} - \frac{(11,5)}{3}\right)^2} = \sqrt{((1.33, 1.33) - (3.67, 1.67))^2}$$
$$= \sqrt{(1.33 - 3.67)^2 + (1.33 - 1.67)^2} = 5.48$$

- $$D_1 = |C_1 - C_2| = \sum_{i=1}^d |C_{1i} - C_{2i}|$$
$$= |1.33 - 3.67| + |1.33 - 1.67| = 2.68$$




# Advantage of CF

- A CF vector provides summary information about the objects a in cluster
- This summary information is sufficient for calculating all the measurements that are needed in BIRCH
- This makes the algorithm efficient because it stores much less data

# Clustering Feature Tree (CF-Tree)

- Height balanced with parameters **B** (branching factor), **L** (leaf entries), **T** (threshold)
- Each **non-leaf node** has at most **B** entries
  - each entry is a CF tuple and a child node link
  - a non-leaf node represents a cluster made of the clusters represented by its entries
- Each **leaf node** has at most **L** CF entries
  - each leaf entry satisfies threshold **T**
  - **T** is maximum diameter (or radius) of any CF in a leaf node
  - CF represents a subcluster
  - each leaf node has links to the next and previous leaf nodes
- Each node must fit a memory page



# Example of CF Tree

Entry:  $[CF_i, child_i]$

Root

$CF_1$	$CF_2$	$CF_3$		$CF_6$
$child_1$	$child_2$	$child_3$		$child_6$

$B = 7$

$L = 6$

Non-leaf node

Entry:  $[CF_i, child_i]$

$CF_1$	$CF_2$	$CF_3$		$CF_5$
$child_1$	$child_2$	$child_3$	.....	$child_5$


Leaf node

Entry:  $[CF_i]$

prev	$CF_1$	$CF_2$	.....	$CF_5$	next
------	--------	--------	-------	--------	------

Leaf node

prev	$CF_1$	$CF_2$	.....	$CF_4$	next
------	--------	--------	-------	--------	------






# CF Tree

- The CF is a very compact representation of the dataset as each entry in a leaf node is not a single data point but a subcluster/cluster
- Tree size is a function of  $T$ 
  - larger  $T$ , more points in each cluster, smaller tree
  - If not enough memory for a given tree, adjust  $T$
- CF tree built dynamically as data is scanned from the dataset and inserted


# Insertion into a CF Tree

- Identify the appropriate leaf
  - start at root, use appropriate distance measure to compare with CF features
  - Recursively follow closest child until you reach a leaf node
- Modify the leaf
  - find closest leaf entry,  $L_i$ , and test if it can absorb new object without violating threshold condition  $T$
  - if yes, update  $CF_i$  for  $L_i$
  - if no and there is space on this leaf for a new entry to hold new object, do so
  - leaves have a max size; may need to be split



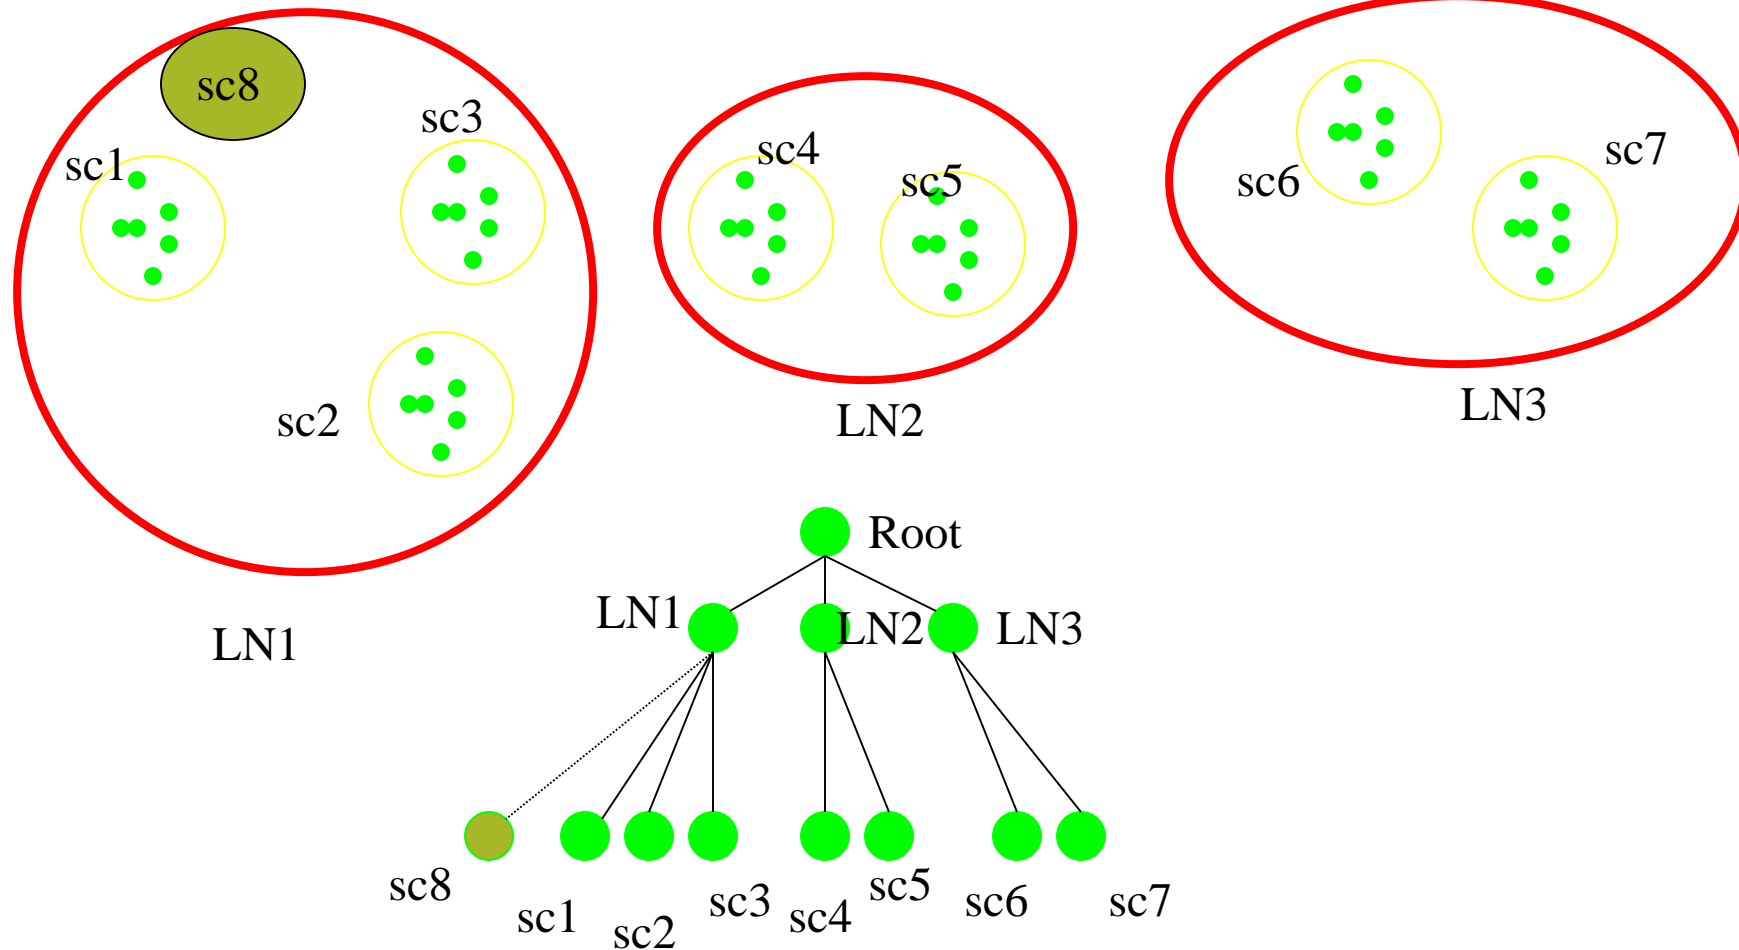
# Insertion into a CF Tree (cont.)

- Modifying the path to the leaf:
  - once object has been inserted, we must update the CF of all ancestors
  - if no node splitting, we only need to update CF vectors
  - a node split requires inserting a new entry into the parent node
  - if parent has space for this entry, we only need to update CF vectors
  - in general, we may need to split the parent (so we do not violate branching factor B) and so on up to the root
  - if root is split, the tree height increases by one

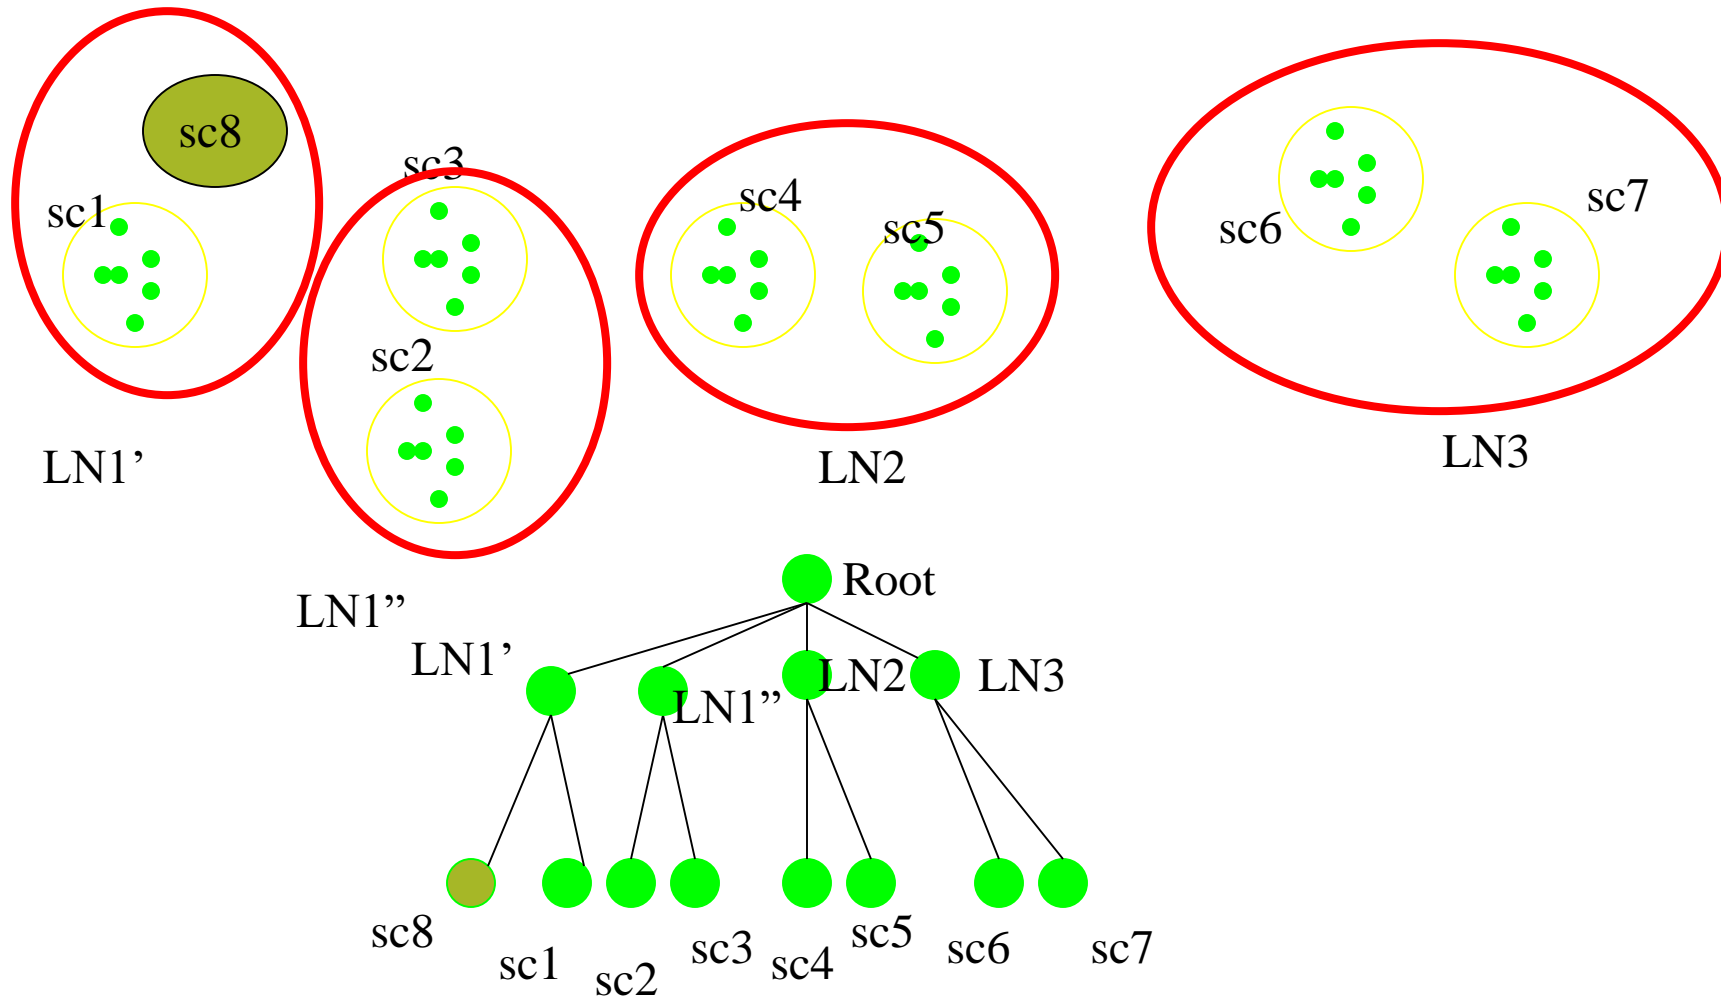


# Example of Insertion - Assume that the sub-clusters are numbered according to the order of formation ( $B=L=3$ )

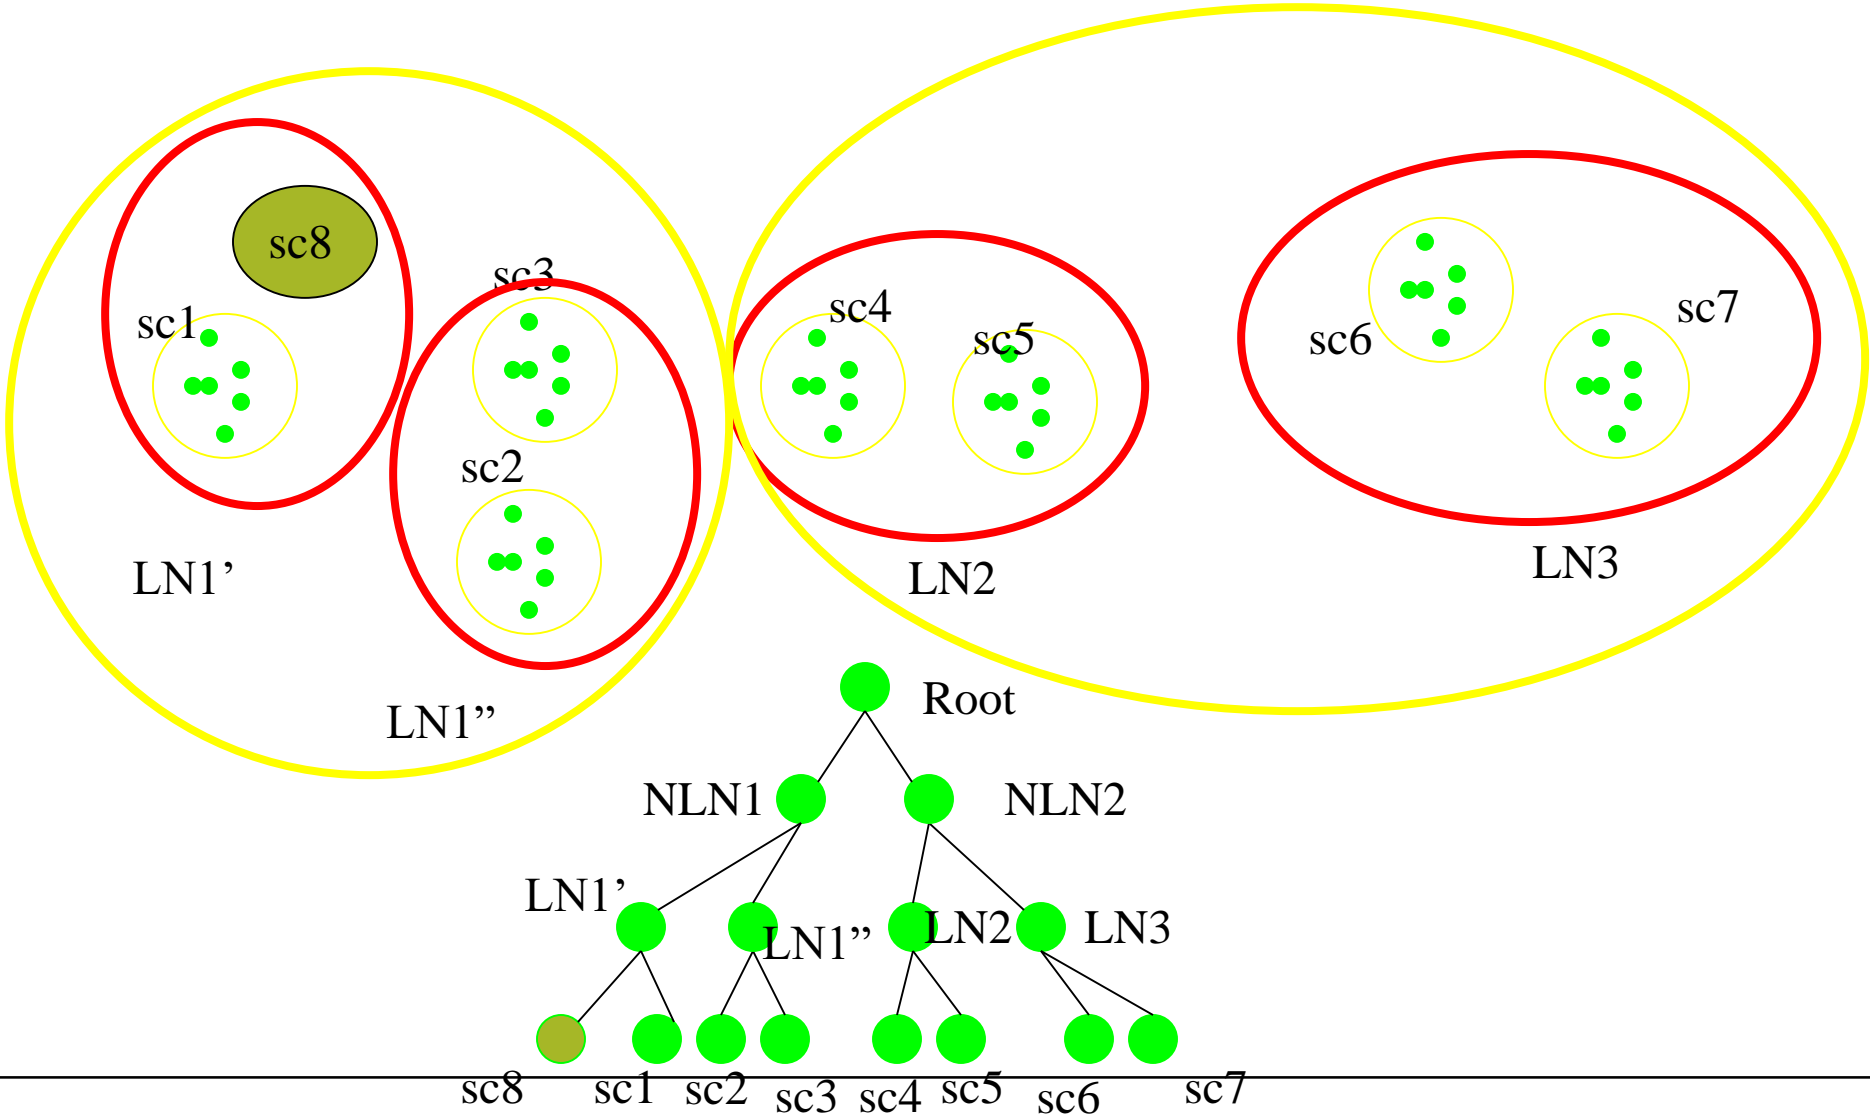
New sub-cluster



If the number of entries of a leaf node (L) can not exceed 3, then LN1 is split.



If the branching factor of a non-leaf node (B) can not exceed 3, then the root is split and the height of the CF Tree increases by one.



# Improving the Clusters

- A node in a CF tree can hold a limited number of entries
- Entries can hold limited number of objects
  - If a data point is inserted twice at different times, the two copies may end up in separate clusters
  - Two sub-clusters that should be in one cluster can be split in different nodes
  - Two sub-clusters that should not be in one cluster can end up in same node
- So, a node may not always correspond to a natural cluster
- These problems can be solved by an additional scan of the dataset, which results in better clustering

# BIRCH Algorithm

- Phase 1 – Build the CF tree
  - If there is not enough memory, increase T and build a smaller CF tree
- Phase 2 – Apply any clustering algorithm to cluster the leaf nodes
- Optional phase to improve the clustering accuracy
  - Scan the whole dataset to re-cluster all data points
  - A data point is placed in the cluster with the closest centroid



# Remarks

- Optional scan solves previously mentioned problems
- BIRCH uses the following ideas to detect outliers:
  - points that are far from any centroid
  - sparse clusters
- BIRCH can handle only numeric data
- BIRCH is a scalable algorithm
- BIRCH does not perform well if clusters are not spherical in shape