

CLUSTERING

CURE and ROCK

CURE

- CURE \equiv Clustering Using REpresentatives
- Problem with many clustering methods
 - favor clusters with spherical “or convex” shapes and similar sizes
 - sensitive to outliers
- CURE can find clusters of any shape and size and it detects outliers
- CURE is a mixed partitional and hierarchical method


Idea

- Several points are used to represent a cluster
- This is more descriptive of clusters of arbitrary shapes
- A set of c well-scattered points from the cluster are used as representatives
 - 1st representative point is farthest from the centroid
 - other scattered points are chosen so that their distance from previously chosen scattered points is maximal


Idea (cont.)

- The well-scattered representative points are then shrunk towards the centroid by a shrinking factor, α
 - When α is 1, all representative points are shrunk to one point, the cluster's centroid
 - When α is 0, no shrinking is done
 - a value between 0.2 to 0.7 is usually used
- Shrinking the scattered points toward the centroid reduces the effects of outliers
- In the hierarchical step, the two closest clusters are merged
 - distance between two clusters is the closest distance between a pair of representatives
- Every time two clusters are merged their representatives are re-calculated
- CURE uses a random sample and partitioning to handle large datasets


CURE Approach




a) Sample of Data



b) Three Clusters with Representative Points



c) Merge Clusters with Closest Points



d) Shrink Representative Points

Outline of CURE

1. Draw a random sample, S , of objects
 - sample size $\sim 2.5\%$ of the size of the dataset
2. Partition the sample S into a set of partitions
 - number of partitions greater than two or three times # of clusters
3. Partially cluster each partition
 - a method like nearest neighbor can be used
 - stop when number of clusters \sim one third number of objects in partition
4. Eliminate outliers
 - outliers are assumed to be small clusters with 1 or 2 objects

Outline of CURE (cont.)

5. Cluster the partial clusters

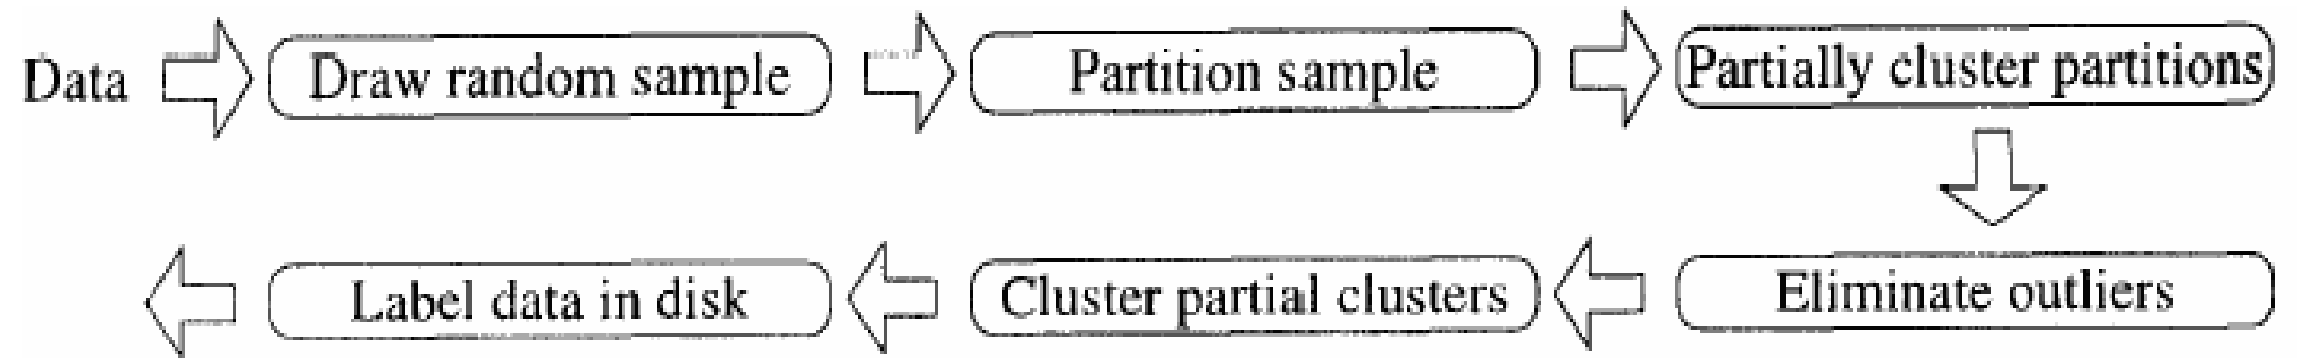
- a hierarchical method is used
- two clusters k_1 and k_2 are merged if they are closest to each other
 $\text{dist}(k_1, k_2) = \min(p, q)$ where $p \in k_1.\text{rep}$, $q \in k_2.\text{rep}$
- find representative points for the new cluster
- shrink representatives by a shrinking factor α , $0 \leq \alpha \leq 1$
- stop when k clusters are found

6. Cluster data on disk

- Assign a point to the cluster with the closest representative


7. Remove small clusters as outliers

Overview of CURE




Experimental Results

- Compared three algorithms on two datasets
 - CURE
 - BIRCH
 - MST (Minimum Spanning Tree)
- Dataset 1 consists of one big circle, two small circles, and two ellipsoids
- Dataset 2 consists of 100 clusters with centers arranged in a grid pattern and points in each cluster following a normal distribution with mean at the cluster center




(a) Data set 1




(b) Data set 2

Experimental Results – Quality of Clustering


- BIRCH split one cluster and merged two clusters
- MST merges the two ellipsoids (long chain effect)
- CURE successfully discovered the clusters in Data set 1



(a) BIRCH



(b) MST METHOD




(c) CURE


Experimental Results – Sensitivity to Parameters

Shrinking factor α :


- 0.2 – 0.7 is a good range of values for α .



(a) $\alpha = 0.1$



(b) $\alpha = 0.2 - 0.7$




(c) $\alpha = 0.8 - 0.9$


Experimental Results – Sensitivity to Parameters (cont.)

Number of Representative Points c :


- For smaller values of c , the quality of clustering suffered.
- However, for values of c greater than 10, CURE always found right clusters.



(a) $c = 2$



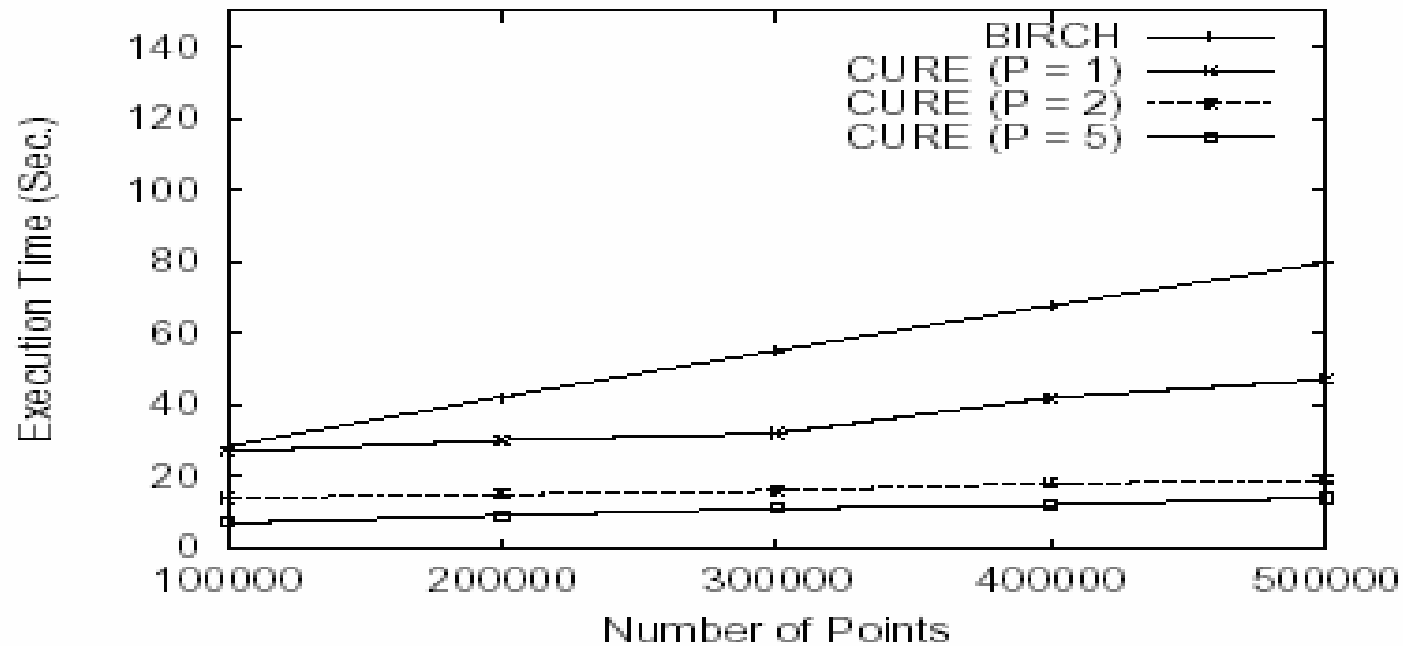
(b) $c = 5$



(c) $c = 10$

Experimental Results – Comparison of Execution time to BIRCH

Run both BIRCH and CURE on Data set 2



Conclusion

- CURE can adjust well to clusters having non-spherical shapes and wide variances in size
- CURE can handle large databases efficiently (one scan)

Clustering with Categorical Attributes

- Algorithms suited for numeric data, do not work well for categorical attributes
- Viewing attribute values as 0/1 pairs or doing simple numeric mapping does not work (next slide)
- Example next page shows that using centroid-based method and Euclidean distance for clustering fails
- There are many clustering algorithms for categorical attributes including: STIRR, ROCK, CACTUS, and c-modes
- We'll talk about the ROCK algorithms

Shortcomings of Traditional Methods on Categorical Data

- Consider a market basket dataset containing 4 transactions over the set of items {a, b, c, d, e, f}

$t_1 = \{a, b, c, e\}$

$t_2 = \{b, c, d, e\}$

$t_3 = \{a, d\}$

$t_4 = \{f\}$

- Represent transactions as 0/1 vectors

- Representation: a b c d e f

$t_1 = (1, 1, 1, 0, 1, 0)$

$t_2 = (0, 1, 1, 1, 1, 0)$

$t_3 = (1, 0, 0, 1, 0, 0)$

$t_4 = (0, 0, 0, 0, 0, 1)$

Shortcomings of Traditional Methods on Categorical Data (cont)

- a b c d e f
 $t_1 = (1, 1, 1, 0, 1, 0)$
 $t_2 = (0, 1, 1, 1, 1, 0)$
 $t_3 = (1, 0, 0, 1, 0, 0)$
 $t_4 = (0, 0, 0, 0, 0, 1)$
- Use centroid-based hierarchical clustering and Euclidean distance to measure distances:
- t_1, t_2 are closest, so merge in cluster, K_1
- centroid of $K_1 = (0.5, 1, 1, 0.5, 1, 0)$
- $d(t_3, K_1) = \sqrt{3.5}$, $d(t_4, k_1) = \sqrt{4.5}$, $d(t_3, t_4) = \sqrt{3}$
- next, algorithm merges t_3 and t_4
- this corresponds to merging transactions $t_3 = \{a, d\}$ and $t_4 = \{f\}$
- so, this method fails on categorical data

	t1	t2	t3	t4
t1	0	$\sqrt{2}$	2	$\sqrt{5}$
t2		0	2	$\sqrt{5}$
t3			0	$\sqrt{3}$
t4				0

ROCK Algorithm

- **ROCK** \equiv **RO**bust **C**lustering using **linK**s
- It is an agglomerative hierarchical clustering method for categorical attributes
- It uses the idea of neighbors and number of links between two items for similarities
- Two objects are **neighbors** if they are similar enough to each other
- **Link** for pair of objects is the number of common neighbors.

Neighbors

- Def: A pair of tuples t_i and t_j are **neighbors** if their similarity is at least some threshold value, θ

i.e., t_i and t_j are neighbors if $\text{sim}(t_i, t_j) \geq \theta$

- The similarity measure used in ROCK is based on Jaccard coefficient and is defined as follows:

$$\text{sim}(t_1, t_2) = \frac{|t_1 \cap t_2|}{|t_1 \cup t_2|}$$

- Example: Find the similarity between $t_1 = \{a, b, c, e\}$ and $t_2 = \{b, c, d, e\}$

$$\text{sim}(t_1, t_2) = \frac{|t_1 \cap t_2|}{|t_1 \cup t_2|} = \frac{|\{a, b, c, e\} \cap \{b, c, d, e\}|}{|\{a, b, c, e\} \cup \{b, c, d, e\}|} = \frac{|\{b, c, e\}|}{|\{a, b, c, d, e\}|} = \frac{3}{5} = 0.6$$

Links

- Def: The **number of links** between two objects t_i and t_j , **$\text{link}(t_i, t_j)$** , is defined as the number of common neighbors they have
- Note that a point is considered as a neighbor of itself

Example

- Consider an information retrieval system where documents have the keywords {Book, Water, Sun, Sand, Swim, Read}
- Given the four documents
d1 = {Book}, **d2** = {Water, Sun, Sand, Swim}
d3 = {Water, Sun, Swim, Read}, **d4** = {Read, Sand}
- similarity matrix using $\text{sim}(t_1, t_2) = \frac{|t_1 \cap t_2|}{|t_1 \cup t_2|}$
- Using threshold $\theta = 0.2$, neighbors for each element are

	d1	d2	d3	d4
d1	1	0	0	0
d2		1	0.6	0.2
d3			1	0.2
d4				1

document	neighbors
d1	d1
d2	d2, d3, d4
d3	d2, d3, d4
d4	d2, d3, d4

	d1	d2	d3	d4
d1	1	0	0	0
d2		3	3	3
d3			3	3
d4				3

- Table showing $\text{link}(d_i, d_j)$
- Resulting Clusters after applying ROCK {{d1}, {d2, d3, d4}}

Outline of the ROCK Algorithm

- Draw a random sample, S , of the data
- Perform clustering on the sample using the link agglomerative hierarchical clustering
 - merge clusters with the largest number of links until k clusters are formed
 - algorithm uses goodness measure $g(C_i, C_j)$ that takes cluster size into consideration
- Cluster data on disk
 - an object is assigned to the cluster with which it has the highest number of links

Remarks

- For a pair of clusters C_i, C_j , **link(C_i, C_j) is the number of mutual links between C_i and C_j**

i.e., $\text{link}(C_i, C_j) = \sum_{t_i \in C_i, t_j \in C_j} \text{link}(t_i, t_j)$

- The goodness measure $g(C_i, C_j)$ is directly proportional to $\text{link}(C_i, C_j)$ and takes into consideration the size of the clusters
- A pair of clusters with maximum g is the best pair to merge