

Bayesian Classification

Bayesian Classification

- A statistical classifier: predicts class membership probabilities
- Foundation: Based on Bayes' Theorem
- Performance: A simple Bayesian classifier, naïve Bayesian classifier, has comparable performance with decision tree and selected neural network classifiers

Bayes' Theorem: Basics

- Let \mathbf{X} be a data sample (“*evidence*”), its class label is unknown
- Let H be a *hypothesis* that X belongs to class C
- Classification is to determine $P(C|\mathbf{X})$, the probability that the hypothesis holds given the observed data sample \mathbf{X}
- $P(H|\mathbf{X})$ or $P(C|\mathbf{X})$ is a *posterior probability*
- $P(C)$ is a *prior probability*
 - the initial probability of C
 - e.g., \mathbf{X} will buy computer, regardless of age, income, ...
- $P(\mathbf{X})$: probability that sample data is observed
- $P(\mathbf{X}|C)$ (likelihood), the probability of observing the sample \mathbf{X} , given that the hypothesis H holds

Bayes' Theorem

- Given training data X , posteriori probability of a hypothesis H , $P(C|X)$, follows the Bayes theorem

$$P(C|X) = \frac{P(X|C)P(C)}{P(X)}$$

- Predicts X belongs to C_i iff the probability $P(C_i|X)$ is the highest among all the $P(C_k|X)$ for all the k classes

Towards Naïve Bayesian Classifier

- Let D be a training dataset
- Each tuple is represented as vector $\mathbf{X} = (x_1, x_2, \dots, x_n)$
- Suppose there are m classes C_1, C_2, \dots, C_m .
- Classification is to find the maximum $P(C_i|\mathbf{X})$
- This can be derived from Bayes' theorem ($1 \leq i \leq m$)

$$P(C_i|\mathbf{X}) = \frac{P(\mathbf{X}|C_i)P(C_i)}{P(\mathbf{X})}$$

- Since $P(\mathbf{X})$ is constant for all classes, only

$$P(C_i|\mathbf{X}) \approx P(\mathbf{X}|C_i)P(C_i)$$

needs to be maximized

Derivation of Naïve Bayes' Classifier

- A simplified assumption: attributes are conditionally independent

$$P(\mathbf{X} | C_i) = \prod_{k=1}^n P(x_k | C_i) = P(x_1 | C_i) \times P(x_2 | C_i) \times \dots \times P(x_n | C_i)$$

- This greatly reduces the computation cost
- If A_k is categorical, $P(x_k | C_i)$ is the # of tuples in C_i having value x_k for A_k divided by $|C_{i,D}|$ (# of tuples of C_i in D)
- If A_k is continuous-valued, $P(x_k | C_i)$ is usually computed based on Gaussian distribution with a mean μ and standard deviation σ

$$g(x, \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

and $P(x_k | C_i)$ is

$$P(\mathbf{X}_k | C_i) = g(x_k, m_{C_i}, S_{C_i})$$

How to Estimate Probabilities from Data?

<i>Tid</i>	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

- Normal distribution:

$$P(A_i|c_j) = \frac{1}{\sqrt{2\pi\sigma_j^2}} e^{-\frac{(A_i-\mu_j)^2}{2\sigma_j^2}}$$

- For (Income, Class=No):
 - If Class=No
 - sample mean = 110K
 - sample variance = 2975K

$$P(\mathbf{X}_k | C_i) = g(x_k, m_{C_i}, S_{C_i})$$

$$P(\text{Income} = 120 | \text{No}) = \frac{1}{\sqrt{2\pi}(54.54)} e^{-\frac{(120-110)^2}{2(2975)}} = 0.0072$$

Naïve Bayes' – Example

Class:

C1:buys_computer = 'yes'

C2:buys_computer = 'no'

Data sample

X = (age <=30,

income = medium,

student = yes

credit_rating = Fair)

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

$$P(C_i|\mathbf{X}) \approx P(\mathbf{X}|C_i)P(C_i)$$

$$P(\mathbf{X}|C_i) = \prod_{k=1}^n P(x_k | C_i) = P(x_1 | C_i) \times P(x_2 | C_i) \times \dots \times P(x_n | C_i)$$

AVC Sets for the Training Dataset

Training Examples

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

AVC-set on *Age*

Age	Buy_Computer	
	yes	no
<=30	2	3
31..40	4	0
>40	3	2

AVC-set on *income*

income	Buy_Computer	
	yes	no
high	2	2
medium	4	2
low	3	1

AVC-set on *Student*

student	Buy_Computer	
	yes	no
yes	6	1
no	3	4

AVC-set on *credit_rating*

Credit rating	Buy_Computer	
	yes	no
fair	6	2
excellent	3	3

Naïve Bayes' – Example

$$P(C_i|\mathbf{X}) \approx P(\mathbf{X}|C_i)P(C_i)$$

- $P(C_i)$: $P(\text{buys_computer} = \text{"yes"}) = 9/14 = 0.643$
 $P(\text{buys_computer} = \text{"no"}) = 5/14 = 0.357$

- Compute $P(\mathbf{X}|C_i)$ for each class

$$P(\text{age} = \text{"<=30"} \mid \text{buys_computer} = \text{"yes"}) = 2/9 = 0.222$$

$$P(\text{age} = \text{"<= 30"} \mid \text{buys_computer} = \text{"no"}) = 3/5 = 0.6$$

$$P(\text{income} = \text{"medium"} \mid \text{buys_computer} = \text{"yes"}) = 4/9 = 0.444$$

$$P(\text{income} = \text{"medium"} \mid \text{buys_computer} = \text{"no"}) = 2/5 = 0.4$$

Age	Buy_Computer	
	yes	no
<=30	2	3
31..40	4	0
>40	3	2

- $\mathbf{X} = (\text{age} \leq 30, \text{income} = \text{medium}, \text{student} = \text{yes}, \text{credit_rating} = \text{fair})$

income	Buy_Computer	
	yes	no
high	2	2
medium	4	2
low	3	1

$$P(\mathbf{X}|C_i) = \prod_{k=1}^n P(x_k | C_i) = P(x_1 | C_i) \times P(x_2 | C_i) \times \dots \times P(x_n | C_i)$$

Naïve Bayes' – Example

$$P(C_i|\mathbf{X}) \approx P(\mathbf{X}|C_i)P(C_i)$$

- $P(C_i)$: $P(\text{buys_computer} = \text{"yes"}) = 9/14 = 0.643$
 $P(\text{buys_computer} = \text{"no"}) = 5/14 = 0.357$

student	Buy_Computer	
	yes	no
yes	6	1
no	3	4

- Compute $P(\mathbf{X}|C_i)$ for each class

$$P(\text{age} = \text{"<=30"} \mid \text{buys_computer} = \text{"yes"}) = 2/9 = 0.222$$

$$P(\text{age} = \text{"<= 30"} \mid \text{buys_computer} = \text{"no"}) = 3/5 = 0.6$$

$$P(\text{income} = \text{"medium"} \mid \text{buys_computer} = \text{"yes"}) = 4/9 = 0.444$$

$$P(\text{income} = \text{"medium"} \mid \text{buys_computer} = \text{"no"}) = 2/5 = 0.4$$

$$P(\text{student} = \text{"yes"} \mid \text{buys_computer} = \text{"yes"}) = 6/9 = 0.667$$

$$P(\text{student} = \text{"yes"} \mid \text{buys_computer} = \text{"no"}) = 1/5 = 0.2$$

$$P(\text{credit_rating} = \text{"fair"} \mid \text{buys_computer} = \text{"yes"}) = 6/9 = 0.667$$

$$P(\text{credit_rating} = \text{"fair"} \mid \text{buys_computer} = \text{"no"}) = 2/5 = 0.4$$

Credit rating	Buy_Computer	
	yes	no
fair	6	2
excellent	3	3

- $\mathbf{X} = (\text{age} \leq 30, \text{income} = \text{medium}, \text{student} = \text{yes}, \text{credit_rating} = \text{fair})$

$$P(\mathbf{X}|C_i) = \prod_{k=1}^n P(x_k | C_i) = P(x_1 | C_i) \times P(x_2 | C_i) \times \dots \times P(x_n | C_i)$$

Naïve Bayes' – Example

$$P(C_i|\mathbf{X}) \approx P(\mathbf{X}|C_i)P(C_i)$$

- $P(C_i)$: $P(\text{buys_computer} = \text{"yes"}) = 9/14 = 0.643$
 $P(\text{buys_computer} = \text{"no"}) = 5/14 = 0.357$
- Compute $P(\mathbf{X}|C_i)$ for each class
 $P(\text{age} = \text{"<=30"} \mid \text{buys_computer} = \text{"yes"}) = 2/9 = 0.222$
 $P(\text{age} = \text{"<= 30"} \mid \text{buys_computer} = \text{"no"}) = 3/5 = 0.6$
 $P(\text{income} = \text{"medium"} \mid \text{buys_computer} = \text{"yes"}) = 4/9 = 0.444$
 $P(\text{income} = \text{"medium"} \mid \text{buys_computer} = \text{"no"}) = 2/5 = 0.4$
 $P(\text{student} = \text{"yes"} \mid \text{buys_computer} = \text{"yes"}) = 6/9 = 0.667$
 $P(\text{student} = \text{"yes"} \mid \text{buys_computer} = \text{"no"}) = 1/5 = 0.2$
 $P(\text{credit_rating} = \text{"fair"} \mid \text{buys_computer} = \text{"yes"}) = 6/9 = 0.667$
 $P(\text{credit_rating} = \text{"fair"} \mid \text{buys_computer} = \text{"no"}) = 2/5 = 0.4$
- $\mathbf{X} = (\text{age} \leq 30, \text{income} = \text{medium}, \text{student} = \text{yes}, \text{credit_rating} = \text{fair})$
 $P(\mathbf{X}|C_i) : P(\mathbf{X} \mid \text{buys_computer} = \text{"yes"}) = 0.222 \times 0.444 \times 0.667 \times 0.667 = 0.044$
 $P(\mathbf{X} \mid \text{buys_computer} = \text{"no"}) = 0.6 \times 0.4 \times 0.2 \times 0.4 = 0.019$
 $P(\mathbf{X}|C_i) * P(C_i) : P(\mathbf{X} \mid \text{buys_computer} = \text{"yes"}) * P(\text{buys_computer} = \text{"yes"}) = 0.028$
 $P(\mathbf{X} \mid \text{buys_computer} = \text{"no"}) * P(\text{buys_computer} = \text{"no"}) = 0.007$

Therefore, \mathbf{X} belongs to class ("buys_computer = yes")

$$P(\mathbf{X} | C_i) = \prod_{k=1}^n P(x_k | C_i) = P(x_1 | C_i) \times P(x_2 | C_i) \times \dots \times P(x_n | C_i)$$

Exercise

- Consider the training dataset given in the table below, where A1, A2 and A3 are the input attributes and C is the class label. Use Naïve Bayes' to find the class for the object $X = \{1, 2, 2\}$.

Sample	Attribute A1	Attribute A2	Attribute A3	Class C
1	1	2	1	1
2	0	0	1	1
3	0	1	2	1
4	1	0	1	1
5	2	1	2	2
6	1	2	1	2
7	2	2	2	2

Avoiding the 0-Probability Problem

- Naïve Bayesian requires each conditional prob. be non-zero. Otherwise, the predicted prob. will be zero

$$P(X | C_i) = \prod_{k=1}^n P(x_k | C_i)$$

- Ex. Suppose that a dataset with 1000 tuples. We have income=low (0), income= medium (990), and income = high (10),
- Use [Laplacian correction](#) (or [Laplacian estimator](#))
 - Adding 1 to each case
 - Prob(income = low) = 1/1003
 - Prob(income = medium) = 991/1003
 - Prob(income = high) = 11/1003
 - The “corrected” prob. estimates are close to their “uncorrected” counterparts

Naïve Bayes' – Advantages and Disadvantages

- Advantages:
 - Easy to implement
 - Very efficient
 - Good results obtained in many applications
 - Can easily handle missing data by omitting that probability in the calculations
- Disadvantages
 - Assumption: class conditional independence, therefore loss of accuracy when the assumption is seriously violated

Naïve Bayes in Scikit-Learn

- There are 3 classes that implement Naïve Bayes in the submodule `sklearn.naive_bayes`
 - GaussianNB
 - MultinomialNB
 - BernoulliNB
- Each assumes a different likelihood distribution of the attributes' values
- Which one to use depends on features' types: continuous, categorical, binary

Preparing the Dataset

```
# Load data
iris = datasets.load_iris()
X = iris.data
y = iris.target

# Split the dataset:
from sklearn.model_selection import
                                train_test_split

X_train, X_test, y_train, y_test =
    train_test_split(X, y, test_size=0.25,
                    random_state=1)
```

Create a GaussianNB Model and Train It

- **`from sklearn.naive_bayes import GaussianNB`**

```
# Create Gaussian Naive Bayes object
```

```
classifier = GaussianNB()
```

```
# Train model
```

```
model = classifier.fit(X_train, y_train)
```

Test the Model

```
# Test the model

y_pred = model.predict(X_test)

from sklearn.metrics import accuracy_score

accuracy_score(y_test, y_pred)
#0.97

new_sample = [[ 5, 4, 3, 2]] # Create new sample

model.predict(new_sample) # Classify it
#array([1])
```

Using NB to Work with Categorical Data

- MultinomialNB is commonly used when working with categorical attributes
 - Ex: movie ratings ranging from 1 to 5
 - Ex: working with text data (e.g., documents of text)
 - Common approach: create a bag of words with a vector for each document containing counts of the appearance of the words in the vocabulary

Prepare the Dataset

```
# Load libraries
```

```
import numpy as np
```

```
from sklearn.naive_bayes import MultinomialNB
```

```
from sklearn.feature_extraction.text import CountVectorizer
```

```
# Create text
```

```
text_data = np.array(['I love Brazil. Brazil!', 'Brazil is best',  
                      'Germany beats both'])
```

```
# Create bag of words
```

```
count = CountVectorizer()
```

```
bag_of_words = count.fit_transform(text_data)
```

CountVectorizer and BOW

Show feature matrix

bag_of_words

Output: <3x7 sparse matrix of type '<class 'numpy.int64'>'
with 8 stored elements in Compressed Sparse Row format>

bag_of_words.toarray()

```
array([[0, 0, 0, 2, 0, 0, 1,],  
       [0, 1, 0, 1, 0, 1, 0],  
       [1, 0, 1, 0, 1, 0, 0]], dtype=int64)
```

Show feature names

count.get_feature_names()

```
['beats', 'best', 'both', 'brazil', 'germany', 'is', 'love']
```

CountVectorizer and BOW (cont.)

- The feature matrix with the words as column names and each row is one observation

beats	best	both	brazil	germany	is	love
0	0	0	2	0	0	1
0	1	0	1	0	1	0
1	0	1	0	1	0	0

- The tuples: (['I love Brazil. Brazil! ', 'Brazil is best', 'Germany beats both'])
- `count.get_feature_names()`
['beats', 'best', 'both', 'brazil', 'germany', 'is', 'love']

Preparing the Dataset and Creating the Model

```
# Create text
```

```
text_data = np.array(['I love Brazil. Brazil!', 'Brazil is best', 'Germany beats both'])
```

```
# Create bag of words
```

```
count = CountVectorizer()
```

```
bag_of_words = count.fit_transform(text_data)
```

```
# Create feature matrix
```

```
features = bag_of_words.toarray()
```

```
# Create target vector
```

```
target = np.array([0,0,1])
```

```
# Create multinomial naive Bayes object with prior probabilities of each class
```

```
classifier = MultinomialNB(class_prior=[0.25, 0.5])
```

```
# Train model
```

```
model = classifier.fit(features, target)
```


Using the Model

Create new observation

```
new_observation = [[0, 0, 0, 1, 0, 1, 0]]
```

Predict new observation's class

```
model.predict(new_observation)
```

```
# array([0])
```

BernoulliNB

- Works similar to MultinomialNB but use with binary data
- See example in the notebook

Distance-Based Classification Algorithms

Distance-Based Classification Algorithms

- Tuples in the same class are more similar to each other
- A similarity or distance measure is used to determine how similar tuples are
- The problem is how to define similarity measures

Distance Measures

- $x_i = (x_{i1}, x_{i2}, \dots, x_{ip})$ and $x_j = (x_{j1}, x_{j2}, \dots, x_{jp})$ are two p -dimensional data objects

- **Minkowski distance:**

$$d(x_i, x_j) = \sqrt[q]{(|x_{i1} - x_{j1}|^q + |x_{i2} - x_{j2}|^q + \dots + |x_{ip} - x_{jp}|^q)}$$

where q is a positive integer

- If $q = 2$, d is **Euclidean distance**

$$d(x_i, x_j) = \sqrt{\sum_{k=1}^p (x_{ik} - x_{jk})^2}$$

- If $q = 1$, d is **Manhattan distance**

$$d(x_i, x_j) = |x_{i1} - x_{j1}| + |x_{i2} - x_{j2}| + \dots + |x_{ip} - x_{jp}|$$

Distance-Based Classification Algorithms


- Place items in class to which they are “closest”
- Must determine distance between an item and a class
- Classes represented by
 - **Centroid**: central value
 - **Medoid**: actual point near the centroid
 - **Individual points**
- Example: $C = \{1, 4, 9\}$
 - Centroid is _____
 - Medoid is _____

Simple Distance Based Classification Algorithm


- Decide on the representative for each class
- Decide on the similarity “or distance” measure to use
- To classify a new sample X , it will be compared to the representative of each class
 - Place X in the same class with the representative that is most similar “or closest” to it

K Nearest Neighbors Algorithm


- Requires three things
 - Set of stored objects
 - Distance measure to compute distance between objects
 - The value of k , the number of nearest neighbors to retrieve
- To classify an unknown object:
 - Compute distance to the training objects
 - Identify k nearest neighbors
 - Use class labels of nearest neighbors to determine the class label of unknown object (e.g., by taking majority vote)




Definition of Nearest Neighbors



(a) 1-nearest neighbor



(b) 2-nearest neighbor



(c) 3-nearest neighbor

K-nearest neighbors of an object x are data points that have the k smallest distance to x

K Nearest Neighbors Algorithm


- Compute distance between two points:
 - Euclidean distance

$$d(p, q) = \sqrt{\sum_i (p_i - q_i)^2}$$

- Determine the class from nearest neighbors list
 - take the majority vote of class labels among the k-nearest neighbors ([simple voting](#))
 - Weigh the vote according to distance ([weighted voting](#))
 - weight factor, $w = 1/d^2$

Choosing the Value of K

- If k is too small, sensitive to noise points
- If k is too large, neighborhood may include points from other classes
- k is usually chosen empirically by trying a range of values



Scaling Values

- Attributes may have to be scaled to prevent distance measures from being dominated by one of the attributes
- Example:
 - height of a person may vary from 1.5m to 1.8m
 - weight of a person may vary from 90lb to 300lb
 - income of a person may vary from \$10K to \$1M

Eager vs Lazy Learners

- KNN is a lazy learner
 - It does not build a model from the training data
 - Unlike eager learners such as decision tree induction and Naïve Bayes
 - Classifying unknown objects is relatively expensive

KNN Algorithm

Input:

D //Training data
 k //Number of neighbors
 t //Input tuple to classify

Output:

c //Class to which t is assigned

KNN Algorithm: //Algorithm to classify tuple using KNN

$N = \emptyset$;

//Find set of neighbors, N , for t

foreach d in D do

 if $|N| < k$ then

$N = N \cup d$;

 else if $\exists u$ in N such that $\text{distance}(t, d) < \text{distance}(t, u)$

 replace d by the neighbor in N with the largest distance to t

c = class to which most u in N are classified

Remarks

- If $K = 1$, nearest neighbor algorithm
- Choice of K is important for KNN
- [Dunham] As a rule of thumb,
 $k \leq \sqrt{\text{number of training samples}}$; commercial algorithms often use a default value of 10.
- Researchers have shown that the classification accuracy of KNN can be as accurate as more elaborated methods
- KNN is slow at the classification time
- KNN does not produce an understandable model

KNN in Scikit-Learn

```
# Load libraries
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn import datasets
```

```
# Load data
```

```
iris = datasets.load_iris()
```

```
X = iris.data
```

```
y = iris.target
```


KNN in Scikit-Learn (cont)

```
# Create standardizer
```

```
standardizer = StandardScaler()
```

```
# Standardize features
```

```
X_std = standardizer.fit_transform(X)
```

```
# Train a KNN classifier with 5 neighbors
```

```
knn = KNeighborsClassifier(n_neighbors=5, n_jobs=-1).fit(X_std, y)
```

```
# Create two observations
```

```
new_observations = [[ 0.75, 0.75, 0.75, 0.75], [ 1, 1, 1, 1]]
```

```
# Predict the class of two observations
```

```
knn.predict(new_observations)
```

```
array([1, 2])
```

KNeighborsClassifier Parameters

- Parameters we can set include:
 - `n_neighbors`: value of K (default is 5)
 - `metric`: 'minkowski' (the default) | 'euclidean' | 'manhattan'
 - `p`: integer, (default = 2) Power parameter for the Minkowski metric
 - `weights`: 'uniform' (the default) | 'distance'
- Notebook name: KNN.ipynb